

# Activity\_Course 2 Waze project lab

September 28, 2025

## 1 Waze Project

### Course 2 - Get Started with Python

Welcome to the Waze Project!

Your Waze data analytics team is still in the early stages of their user churn project. Previously, you were asked to complete a project proposal by your supervisor, May Santner. You have received notice that your project proposal has been approved and that your team has been given access to Waze's user data. To get clear insights, the user data must be inspected and prepared for the upcoming process of exploratory data analysis (EDA).

A Python notebook has been prepared to guide you through this project. Answer the questions and create an executive summary for the Waze data team.

## 2 Course 2 End-of-course project: Inspect and analyze data

In this activity, you will examine data provided and prepare it for analysis. This activity will help ensure the information is,

1. Ready to answer questions and yield insights
2. Ready for visualizations
3. Ready for future hypothesis testing and statistical methods

**The purpose** of this project is to investigate and understand the data provided.

**The goal** is to use a dataframe constructed within Python, perform a cursory inspection of the provided dataset, and inform team members of your findings.

*This activity has three parts:*

**Part 1:** Understand the situation \* How can you best prepare to understand and organize the provided information?

**Part 2:** Understand the data

- Create a pandas dataframe for data learning, future exploratory data analysis (EDA), and statistical activities
- Compile summary information about the data to inform next steps

### Part 3: Understand the variables

- Use insights from your examination of the summary data to guide deeper investigation into variables

Follow the instructions and answer the following questions to complete the activity. Then, you will complete an Executive Summary using the questions listed on the PACE Strategy Document.

Be sure to complete this activity before moving on. The next course item will provide you with a completed exemplar to compare to your own work.

## 3 Identify data types and compile summary information

### 4 PACE stages

Throughout these project notebooks, you'll see references to the problem-solving framework, PACE. The following notebook components are labeled with the respective PACE stages: Plan, Analyze, Construct, and Execute.

#### 4.1 PACE: Plan

Consider the questions in your PACE Strategy Document and those below to craft your response:

##### 4.1.1 Task 1. Understand the situation

- How can you best prepare to understand and organize the provided driver data?  
> *Begin by exploring your dataset and consider reviewing the Data Dictionary.*
- 1. **Understand the Business Context:** Clearly define the project's goal (churn prediction) and how churn is defined by Waze.
- 2. **Detailed Data Review:** Beyond basic `head()` and `info()`, thoroughly examine `describe()` output, check unique values for categorical columns, and understand variable relationships.
- 3. **Review and confirm the Data Dictionary:** Confirm the meaning, units, and potential range of values for every column. If anything is unclear or missing from the provided dictionary, that's when we'd add to it or seek clarification.
- What follow-along and self-review codebooks will help you perform this work?  
> 1. **Python Libraries:** Core libraries like Pandas (for data manipulation), NumPy (for numerical ops), and Matplotlib/Seaborn (for visualization).  
2. **Interactive Environments:** Use *Jupyter Notebooks* or *Google Colab*. They allow you to combine code, output, and explanations, making it easy to follow your own logic and review.  
3. **Clear Documentation:** Use *Markdown cells* within your notebooks to explain steps, observations, and decisions. This acts as your personal "codebook."

- What are some additional activities a resourceful learner would perform before starting to code? > \* **Formulate Hypotheses** (optional): Brainstorm questions you want the data to answer (e.g., “Do users who drive fewer days churn more?”).
  - **Sketch Out Analysis Plan**: Outline the high-level steps: data cleaning, initial exploration, feature engineering, modeling, and evaluation.
  - **Research Domain Knowledge** (optional): Learn more about typical Waze usage patterns, factors influencing app churn, and common metrics in mobile analytics.
  - **Consider Data Limitations**: Think about what questions the data *cannot* answer and what biases might exist.

## 4.2 PACE: Analyze

Consider the questions in your PACE Strategy Document to reflect on the Analyze stage.

### 4.2.1 Task 2a. Imports and data loading

Start by importing the packages that you will need to load and explore the dataset. Make sure to use the following import statements:

- `import pandas as pd`
- `import numpy as np`

```
[1]: # Import packages for data manipulation
import pandas as pd
import numpy as np

from IPython.display import display
```

Then, load the dataset into a dataframe. Creating a dataframe will help you conduct data manipulation, exploratory data analysis (EDA), and statistical activities.

**Note:** As shown in this cell, the dataset has been automatically loaded in for you. You do not need to download the .csv file, or provide more code, in order to access the dataset and proceed with this lab. Please continue with this activity by completing the following instructions.

```
[2]: # Load dataset into dataframe
df = pd.read_csv('waze_dataset.csv')
```

### 4.2.2 Task 2b. Summary information

View and inspect summary information about the dataframe by **coding the following**:

1. `df.head(10)`
2. `df.info()`

*Consider the following questions:*

**Question 1:** When reviewing the `df.head()` output, are there any variables that have missing values?

**Question 2:** When reviewing the `df.info()` output, what are the data types? How many rows and columns do you have?

**Question 3:** Does the dataset have any missing values?

```
[3]: # Display and examine the first ten rows of the dataframe
df.head(10)
```

```
[3]:
```

	ID	label	sessions	drives	total_sessions	n_days_after_onboarding	\
0	0	retained	283	226	296.748273		2276
1	1	retained	133	107	326.896596		1225
2	2	retained	114	95	135.522926		2651
3	3	retained	49	40	67.589221		15
4	4	retained	84	68	168.247020		1562
5	5	retained	113	103	279.544437		2637
6	6	retained	3	2	236.725314		360
7	7	retained	39	35	176.072845		2999
8	8	retained	57	46	183.532018		424
9	9	churned	84	68	244.802115		2997

	total_navigations_fav1	total_navigations_fav2	driven_km_drives	\
0	208	0	2628.845068	
1	19	64	13715.920550	
2	0	0	3059.148818	
3	322	7	913.591123	
4	166	5	3950.202008	
5	0	0	901.238699	
6	185	18	5249.172828	
7	0	0	7892.052468	
8	0	26	2651.709764	
9	72	0	6043.460295	

	duration_minutes_drives	activity_days	driving_days	device
0	1985.775061	28	19	Android
1	3160.472914	13	11	iPhone
2	1610.735904	14	8	Android
3	587.196542	7	3	iPhone
4	1219.555924	27	18	Android
5	439.101397	15	11	iPhone
6	726.577205	28	23	iPhone
7	2466.981741	22	20	iPhone
8	1594.342984	25	20	Android
9	2341.838528	7	3	iPhone

```
[4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                    14999 non-null  int64
1   label                                14299 non-null  object
2   sessions                            14999 non-null  int64
3   drives                              14999 non-null  int64
4   total_sessions                      14999 non-null  float64
5   n_days_after_onboarding             14999 non-null  int64
6   total_navigations_fav1              14999 non-null  int64
7   total_navigations_fav2              14999 non-null  int64
8   driven_km_drives                    14999 non-null  float64
9   duration_minutes_drives              14999 non-null  float64
10  activity_days                       14999 non-null  int64
11  driving_days                        14999 non-null  int64
12  device                              14999 non-null  object
dtypes: float64(3), int64(8), object(2)
memory usage: 1.5+ MB

```

```

[5]: summary_df = pd.DataFrame({
    'Column Name': df.columns,
    'Non-Null Count': df.notnull().sum().values,
    'Null Count': df.isnull().sum().values,
    'Data Type' : df.dtypes.values,
    '% Missing' : ((df.isnull().sum().values / len(df)) * 100).round(2)
})
display(summary_df)

```

	Column Name	Non-Null Count	Null Count	Data Type	% Missing
0	ID	14999	0	int64	0.00
1	label	14299	700	object	4.67
2	sessions	14999	0	int64	0.00
3	drives	14999	0	int64	0.00
4	total_sessions	14999	0	float64	0.00
5	n_days_after_onboarding	14999	0	int64	0.00
6	total_navigations_fav1	14999	0	int64	0.00
7	total_navigations_fav2	14999	0	int64	0.00
8	driven_km_drives	14999	0	float64	0.00
9	duration_minutes_drives	14999	0	float64	0.00
10	activity_days	14999	0	int64	0.00
11	driving_days	14999	0	int64	0.00
12	device	14999	0	object	0.00

```

[6]: # Get summary statistics
df.describe()

```

```
[6]:
```

	ID	sessions	drives	total_sessions	\
count	14999.000000	14999.000000	14999.000000	14999.000000	
mean	7499.000000	80.633776	67.281152	189.964447	
std	4329.982679	80.699065	65.913872	136.405128	
min	0.000000	0.000000	0.000000	0.220211	
25%	3749.500000	23.000000	20.000000	90.661156	
50%	7499.000000	56.000000	48.000000	159.568115	
75%	11248.500000	112.000000	93.000000	254.192341	
max	14998.000000	743.000000	596.000000	1216.154633	

	n_days_after_onboarding	total_navigations_fav1	\
count	14999.000000	14999.000000	
mean	1749.837789	121.605974	
std	1008.513876	148.121544	
min	4.000000	0.000000	
25%	878.000000	9.000000	
50%	1741.000000	71.000000	
75%	2623.500000	178.000000	
max	3500.000000	1236.000000	

	total_navigations_fav2	driven_km_drives	duration_minutes_drives	\
count	14999.000000	14999.000000	14999.000000	
mean	29.672512	4039.340921	1860.976012	
std	45.394651	2502.149334	1446.702288	
min	0.000000	60.441250	18.282082	
25%	0.000000	2212.600607	835.996260	
50%	9.000000	3493.858085	1478.249859	
75%	43.000000	5289.861262	2464.362632	
max	415.000000	21183.401890	15851.727160	

	activity_days	driving_days
count	14999.000000	14999.000000
mean	15.537102	12.179879
std	9.004655	7.824036
min	0.000000	0.000000
25%	8.000000	5.000000
50%	16.000000	12.000000
75%	23.000000	19.000000
max	31.000000	30.000000

## RESPONSES TO QUESTIONS 1-3 HERE

**Answer 1:** > Based on a visual inspection of the `df.head()` output provided, there are no immediately apparent missing values. All cells in the first 10 rows appear to be populated.

**Answer 2:** > The dataset contains eight columns with a data type of `int64`, three columns with a data type of `float64`, and two columns with a data type of `object`. The dataset has **14,999 rows** (represents one unique user) and **13 columns**.

**Answer 3:** > Yes, the dataset has missing values. The `df.info()` output shows that the `label` column has only 14,299 non-null entries out of a total of 14,999 rows. This indicates there are **700 missing values** in the `label` column.

## Initial Cleaning

- Replace columnn name `ID` with `id` for consistency
- Converting `label` and `device` to category:
  - Saves memory
  - Speeds up operations like grouping, filtering
  - Makes the meaning of the data clearer (they're not meant for math)
- Converting `device` to string

```
[7]: waze_users_df = df.rename(columns={'ID': 'id'})
waze_users_df['label'] = waze_users_df['label'].astype('category')
waze_users_df['device'] = waze_users_df['device'].astype('category')
waze_users_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     14999 non-null  int64
1   label                                14299 non-null  category
2   sessions                             14999 non-null  int64
3   drives                               14999 non-null  int64
4   total_sessions                       14999 non-null  float64
5   n_days_after_onboarding              14999 non-null  int64
6   total_navigations_fav1               14999 non-null  int64
7   total_navigations_fav2               14999 non-null  int64
8   driven_km_drives                     14999 non-null  float64
9   duration_minutes_drives              14999 non-null  float64
10  activity_days                        14999 non-null  int64
11  driving_days                         14999 non-null  int64
12  device                               14999 non-null  category
dtypes: category(2), float64(3), int64(8)
memory usage: 1.3 MB
```

### 4.2.3 Task 2c. Null values and summary statistics

Compare the summary statistics of the 700 rows that are missing labels with summary statistics of the rows that are not missing any values.

**Question:** Is there a discernible difference between the two populations?

```
[8]: # Isolate rows with null values
null_label_waze_users_df = waze_users_df[waze_users_df['label'].isnull()]
# null_label_waze_users_df.head(10)
```

```
[9]: # Display summary stats of rows with null values
null_label_waze_users_df.describe()
```

```
[9]:
```

	id	sessions	drives	total_sessions	\
count	700.000000	700.000000	700.000000	700.000000	
mean	7405.584286	80.837143	67.798571	198.483348	
std	4306.900234	79.987440	65.271926	140.561715	
min	77.000000	0.000000	0.000000	5.582648	
25%	3744.500000	23.000000	20.000000	94.056340	
50%	7443.000000	56.000000	47.500000	177.255925	
75%	11007.000000	112.250000	94.000000	266.058022	
max	14993.000000	556.000000	445.000000	1076.879741	

	n_days_after_onboarding	total_navigations_fav1	\
count	700.000000	700.000000	
mean	1709.295714	118.717143	
std	1005.306562	156.308140	
min	16.000000	0.000000	
25%	869.000000	4.000000	
50%	1650.500000	62.500000	
75%	2508.750000	169.250000	
max	3498.000000	1096.000000	

	total_navigations_fav2	driven_km_drives	duration_minutes_drives	\
count	700.000000	700.000000	700.000000	
mean	30.371429	3935.967029	1795.123358	
std	46.306984	2443.107121	1419.242246	
min	0.000000	290.119811	66.588493	
25%	0.000000	2119.344818	779.009271	
50%	10.000000	3421.156721	1414.966279	
75%	43.000000	5166.097373	2443.955404	
max	352.000000	15135.391280	9746.253023	

	activity_days	driving_days
count	700.000000	700.000000
mean	15.382857	12.125714
std	8.772714	7.626373
min	0.000000	0.000000
25%	8.000000	6.000000
50%	15.000000	12.000000
75%	23.000000	18.000000
max	31.000000	30.000000



```
[10]: # Isolate rows without null values
non_null_label_waze_users_df = waze_users_df[~waze_users_df['label'].isnull()]
# non_null_label_waze_users_df.head(10)
```

```
[11]: # Display summary stats of rows without null values
non_null_label_waze_users_df.describe()
```

```
[11]:
```

	id	sessions	drives	total_sessions	\
count	14299.000000	14299.000000	14299.000000	14299.000000	
mean	7503.573117	80.623820	67.255822	189.547409	
std	4331.207621	80.736502	65.947295	136.189764	
min	0.000000	0.000000	0.000000	0.220211	
25%	3749.500000	23.000000	20.000000	90.457733	
50%	7504.000000	56.000000	48.000000	158.718571	
75%	11257.500000	111.000000	93.000000	253.540450	
max	14998.000000	743.000000	596.000000	1216.154633	

	n_days_after_onboarding	total_navigations_fav1	\
count	14299.000000	14299.000000	
mean	1751.822505	121.747395	
std	1008.663834	147.713428	
min	4.000000	0.000000	
25%	878.500000	10.000000	
50%	1749.000000	71.000000	
75%	2627.500000	178.000000	
max	3500.000000	1236.000000	

	total_navigations_fav2	driven_km_drives	duration_minutes_drives	\
count	14299.000000	14299.000000	14299.000000	
mean	29.638296	4044.401535	1864.199794	
std	45.350890	2504.977970	1448.005047	
min	0.000000	60.441250	18.282082	
25%	0.000000	2217.319909	840.181344	
50%	9.000000	3496.545617	1479.394387	
75%	43.000000	5299.972162	2466.928876	
max	415.000000	21183.401890	15851.727160	

	activity_days	driving_days
count	14299.000000	14299.000000
mean	15.544653	12.182530
std	9.016088	7.833835
min	0.000000	0.000000
25%	8.000000	5.000000
50%	16.000000	12.000000
75%	23.000000	19.000000
max	31.000000	30.000000

**Answer:** > There is no discernible difference in both the rows with null values in the label column and the rows with non-null values. The mean values are extremely close for both population, there are no significant differences for other variable metrics as well.

This finding suggests that the missing values in the label column may be missing at random, and are not due to any specific user behavior patterns.

#### 4.2.4 Task 2d. Null values - device counts

Next, check the two populations with respect to the `device` variable.

**Question:** How many iPhone users had null values and how many Android users had null values?

```
[12]: # Get count of null values by device
null_label_waze_users_df['device'].value_counts()
```

```
[12]: iPhone      447
      Android    253
      Name: device, dtype: int64
```

**Answer:** > Out of 700 rows with null label value, 447 are iPhone users and 253 are Android users.

Now, of the rows with null values, calculate the percentage with each device—Android and iPhone. You can do this directly with the `value_counts()` function.

```
[13]: # Calculate % of iPhone nulls and Android nulls
null_label_waze_users_df['device'].value_counts(normalize=True).round(4) * 100
```

```
[13]: iPhone      63.86
      Android    36.14
      Name: device, dtype: float64
```

**Question:** How does this compare to the device ratio in the full dataset?

```
[14]: # Calculate % of iPhone users and Android users in full dataset
waze_users_df['device'].value_counts(normalize=True).round(4) * 100
```

```
[14]: iPhone      64.48
      Android    35.52
      Name: device, dtype: float64
```

**Answer:** > The percentage of missing values by each device is consistent with their representation in the data overall.

There is nothing to suggest a non-random cause of the missing data.

Examine the counts and percentages of users who churned vs. those who were retained.

**Question:** How many of each group are represented in the data?

```
[15]: # Calculate counts of churned vs. retained
print(waze_users_df['label'].value_counts())
print('-----')
print(waze_users_df['label'].value_counts(normalize=True).round(4) * 100)
```

```
retained    11763
churned      2536
Name: label, dtype: int64
-----
retained     82.26
churned     17.74
Name: label, dtype: float64
```

**Answer:** > This dataset contains 82% retained users and 18% churned users.

Next, compare the medians of each variable for churned and retained users. The reason for calculating the median and not the mean is that you don't want outliers to unduly affect the portrayal of a typical user. Notice, for example, that the maximum value in the `driven_km_drives` column is 21,183 km. That's more than half the circumference of the earth!

```
[16]: waze_users_label_group = waze_users_df.groupby(['label'])
```

```
[17]: pd.set_option('display.max_columns', None)
display(waze_users_label_group.agg(['mean', 'median', 'min', 'max']))
```

	id				sessions				
	mean	median	min	max	mean	median	min	max	\
label									
churned	7544.852918	7477.5	9	14997	87.238959	59.0	0	743	
retained	7494.673553	7509.0	0	14998	79.197654	56.0	0	725	

	drives				total_sessions			
	mean	median	min	max	mean	median	min	\
label								
churned	72.730678	50.0	0	596	196.893424	164.339042	1.362129	
retained	66.075491	47.0	0	582	187.963672	157.586756	0.220211	

	n_days_after_onboarding					
	max	mean	median	min	max	\
label						
churned	1216.154633	1471.027603	1321.0	6	3496	
retained	1117.893821	1812.359432	1843.0	4	3500	

	total_navigations_fav1				total_navigations_fav2		
	mean	median	min	max	mean		\
label							
churned	139.414826	84.5	0	1170	31.596609		
retained	117.938451	68.0	0	1236	29.216101		

	driven_km_drives				\	
	median	min	max	mean	median	min
label						
churned	11.0	0	396	4147.171864	3652.655666	178.232313
retained	9.0	0	415	4022.245150	3464.684614	60.441250

	duration_minutes_drives				\	
	max	mean	median	min		
label						
churned	19214.47511	1975.459630	1607.183785	23.022685		
retained	21183.40189	1840.213146	1458.046141	18.282082		

	activity_days				driving_days				
	max	mean	median	min	max	mean	median	min	max
label									
churned	10040.56896	9.644716	8.0	0	31	7.218060	6.0	0	29
retained	15851.72716	16.816628	17.0	0	31	13.252827	14.0	0	30

```
[18]: # Calculate median values of all columns for churned and retained users
display(waze_users_label_group.median(numeric_only=True))
```

	id	sessions	drives	total_sessions	n_days_after_onboarding	\
label						
churned	7477.5	59.0	50.0	164.339042		1321.0
retained	7509.0	56.0	47.0	157.586756		1843.0

	total_navigations_fav1	total_navigations_fav2	driven_km_drives	\
label				
churned		84.5	11.0	3652.655666
retained		68.0	9.0	3464.684614

	duration_minutes_drives	activity_days	driving_days
label			
churned	1607.183785	8.0	6.0
retained	1458.046141	17.0	14.0

This offers an interesting snapshot of the two groups, churned vs. retained:

Users who churned averaged ~3 more drives in the last month than retained users, but retained users used the app on over twice as many days as churned users in the same time period.

The median churned user drove ~200 more kilometers and 2.5 more hours during the last month than the median retained user.

It seems that churned users had more drives in fewer days, and their trips were farther and longer in duration. Perhaps this is suggestive of a user profile. Continue exploring!

Calculate the median kilometers per drive in the last month for both retained and churned users.

Begin by dividing the `driven_km_drives` column by the `drives` column. Then, group the results by churned/retained and calculate the median km/drive of each group.

```
[19]: # Add a column to df called `km_per_drive`
waze_users_df['km_per_drive'] = waze_users_df['driven_km_drives'] / \
    ↪waze_users_df['drives']

# Group by `label`, calculate the median, and isolate for km per drive
waze_users_label_group = waze_users_df.groupby(['label'])
waze_users_label_group[['km_per_drive']].median(numeric_only=True)
```

```
[19]:          km_per_drive
label
churned      74.109416
retained     75.014702
```

The median retained user drove about one more kilometer per drive than the median churned user.

**Question:** How many kilometers per driving day was this?

To calculate this statistic, repeat the steps above using `driving_days` instead of `drives`.

```
[20]: # Add a column to df called `km_per_driving_day`
waze_users_df['km_per_driving_day'] = waze_users_df['driven_km_drives'] / \
    ↪waze_users_df['driving_days']

# Group by `label`, calculate the median, and isolate for km per driving day
waze_users_label_group = waze_users_df.groupby(['label'])
waze_users_label_group[['km_per_driving_day']].median(numeric_only=True)
```

```
[20]:          km_per_driving_day
label
churned      697.541999
retained     289.549333
```

Now, calculate the median number of drives per driving day for each group.

```
[21]: # Add a column to df called `drives_per_driving_day`
waze_users_df['drives_per_driving_day'] = waze_users_df['drives'] / \
    ↪waze_users_df['driving_days']

# Group by `label`, calculate the median, and isolate for drives per driving day
waze_users_label_group = waze_users_df.groupby(['label'])
waze_users_label_group[['drives_per_driving_day']].median(numeric_only=True)
```

```
[21]:          drives_per_driving_day
label
churned      10.0000
retained      4.0625
```

**Answer:** > The median user who churned drove 698 kilometers each day they drove last month, which is almost ~240% the per-drive-day distance of retained users. The median churned user had a similarly disproportionate number of drives per drive day compared to retained users.

It is clear from these figures that, regardless of whether a user churned or not, the users represented in this data are serious drivers! It would probably be safe to assume that this data does not represent typical drivers at large. Perhaps the data—and in particular the sample of churned users—contains a high proportion of long-haul truckers.

In consideration of how much these users drive, it would be worthwhile to recommend to Waze that they gather more data on these super-drivers. It's possible that the reason for their driving so much is also the reason why the Waze app does not meet their specific set of needs, which may differ from the needs of a more typical driver, such as a commuter.

Finally, examine whether there is an imbalance in how many users churned by device type.

Begin by getting the overall counts of each device type for each group, churned and retained.

```
[22]: # For each label, calculate the number of Android users and iPhone users
      # waze_users_label_group['device'].value_counts()
      waze_users_label_device_group = waze_users_df.groupby(['label', 'device'])
      waze_users_label_device_group.size()
```

```
[22]: label      device
      churned  Android      891
           iPhone    1645
      retained  Android    4183
           iPhone    7580
      dtype: int64
```

Now, within each group, churned and retained, calculate what percent was Android and what percent was iPhone.

```
[23]: # For each label, calculate the percentage of Android users and iPhone users
      waze_users_label_group['device'].value_counts(normalize=True).round(4) * 100
```

```
[23]: label
      churned  iPhone    64.87
           Android    35.13
      retained  iPhone    64.44
           Android    35.56
      Name: device, dtype: float64
```

The ratio of iPhone users and Android users is consistent between the churned group and the retained group, and those ratios are both consistent with the ratio found in the overall dataset.

---

```
[24]: # Import vizualization packages
      import seaborn as sns
      import matplotlib.pyplot as plt
```

```

from matplotlib.ticker import FuncFormatter
from matplotlib.ticker import PercentFormatter
from matplotlib.cbook import boxplot_stats
from IPython.display import display, HTML

# Formatter to display y-axis numbers with commas
formatter = FuncFormatter(lambda x, _: f'{int(x):,}')

# Set style
sns.set(style='whitegrid')

```

```

[25]: # Helper function to safely add text if value is finite
def safe_text(ax, x, y, text, **kwargs):
    if np.isfinite(y):
        ax.text(x, y, f'{text}:\n{y:.2f}', va='center', fontsize=8, **kwargs)

def summarize_column_stats(source_df, columns_of_interest, bool_mask=None,
    mask_label=None):
    # Apply mask if any
    masked_df = source_df
    if bool_mask is not None:
        masked_df = source_df[bool_mask]

    # Select only the columns of interest
    subset_df = masked_df[columns_of_interest]

    # Combine into one DataFrame
    summary_df = pd.DataFrame({
        'total': len(subset_df), # Total row count (same for
    all columns)
        'non-null': subset_df.count(), # Count non-null values
        'null': subset_df.isna().sum(), # Count null values
        'min': subset_df.min(numeric_only=True), # Minimum value
        'max': subset_df.max(numeric_only=True) # Maximum value
    })

    # If mask label is provided, show it
    if mask_label:
        display(HTML(f"<span>Summary for condition: <u>{mask_label}</u></
    span>"))
    print(summary_df)

#
    -----
# NOT A GENERIC FUNCTION - SPECIFIC TO WAZE DATA

```

```

def summarize_waze_column_stats_multi(columns_of_interest, mask_column_name,
    mask_column_label):
    summarize_column_stats(waze_users_df, columns_of_interest,
        waze_users_df[mask_column_name] == 0,
    f'{mask_column_label} that have zero values')
    summarize_column_stats(waze_users_df, columns_of_interest,
        (waze_users_df[mask_column_name] == 0) &
    (waze_users_df['label'] == 'churned'),
        f'{mask_column_label} that have zero values for
    Churned Users')
    summarize_column_stats(waze_users_df, columns_of_interest,
        (waze_users_df[mask_column_name] == 0) &
    (waze_users_df['label'] == 'retained'),
        f'{mask_column_label} that have zero values for
    Retained Users')
    summarize_column_stats(waze_users_df, columns_of_interest,
        (waze_users_df[mask_column_name] == 0) &
    (waze_users_df['label'].isna()),
        f'{mask_column_label} that have zero values for NULL
    User Labels')
#
    -----

def plot_bar_by_group(source_df,
    x_column_name,
    x_column_label,
    hue_column_name,
    hue_column_label,
    count_label):
    # Grouped data
    grouped_data = source_df.groupby([x_column_name, hue_column_name]).size().
    reset_index(name='count')

    plt.figure(figsize=(10, 6))
    barplot = sns.barplot(data=grouped_data, x=x_column_name, y='count',
    hue=hue_column_name)

    # Manually add text labels
    for bar in barplot.patches:
        height = bar.get_height()
        x = bar.get_x() + bar.get_width() / 2 # Center the text on the bar
        plt.text(x, height + 100, int(height), ha='center', va='bottom',
    fontsize=9)

    plt.title(f'Distribution of {count_label} by {x_column_label} and
    {hue_column_label}')

```



```

plt.xlabel(f'{x_column_label}')
plt.ylabel(f'{count_label}')
plt.legend(title=hue_column_label)
plt.tight_layout()
plt.show()

def group_summary_stats(grouped_df, column_name):
    # Define the aggregation functions
    agg_funcs = {
        column_name: [lambda x: round(x.mean(), 2),
                       'median',
                       lambda x: x.quantile(0.25),
                       lambda x: x.quantile(0.75),
                       lambda x: x.quantile(0.75) - x.quantile(0.25),
                       'min',
                       'max']

    }

    # Perform the aggregation
    stats_df = grouped_df.agg(agg_funcs)
    stats_df.columns = ['mean', 'median', 'Q1', 'Q3', 'IQR', 'min', 'max']

    # Reset index for clean output
    stats_df = stats_df.reset_index()

    pd.set_option('display.max_columns', None)
    return stats_df

def plot_box_by_group(source_df,
                      ax,
                      column_name,
                      column_label,
                      group_by_column_name,
                      group_by_column_label,
                      grouped_df=None,
                      show_legend=False):

    sns.boxplot(data=source_df, x=group_by_column_name, y=column_name, ax=ax)
    ax.yaxis.set_major_formatter(formatter)
    ax.set_title(f'Distribution of {column_label} by {group_by_column_label}')
    ax.set_xlabel(f'{group_by_column_label}')
    ax.set_ylabel(f'{column_label}')

    # Plot means
    means = grouped_df[column_name].mean()
    for i, mean in enumerate(means):
        if np.isfinite(mean): # Safeguard against NaN, inf, -inf

```

```

        ax.scatter(i, mean, color='red', marker='D',
                    s=40, zorder=5, label='Mean' if i == 0 else "")
        ax.hlines(mean, xmin=i - 0.4, xmax=i + 0.4,
                  colors='red', linestyles='dashed', linewidth=1)
        ax.text(i, mean, f'{mean:.2f}',
                ha='center', va='bottom', fontsize=8)

    # Calculate boxplot stats
    stats = []
    for i, (name, group) in enumerate(grouped_df):
        # stats = boxplot_stats(group[column_name])[0] # Get stats for this
        # group

        # Remove NaN values from the column_name
        column_data = group[column_name].dropna()

        # Replace positive/negative infinity with very large/small placeholders
        # so they can still be included in the stats calculation
        temp_data = column_data.replace(np.inf, 1e12).replace(-np.inf, -1e12)

        # Calculate boxplot statistics (min, Q1, median, Q3, max, etc.)
        stats = boxplot_stats(temp_data, whis=1.5)[0]

        # Min and Max (whiskers)
        safe_text(ax, i - 0.2, stats['whislo'], 'Lower Whisker', ha='right')
        safe_text(ax, i - 0.2, stats['whishi'], 'Upper Whisker', ha='right')

        # Q1 and Q3
        safe_text(ax, i + 0.2, stats['q1'], 'Q1', ha='left')
        safe_text(ax, i + 0.2, stats['q3'], 'Q3', ha='left')

    # Show legend only if there are labeled plot elements.
    # If no handles are found, it means all relevant values were NaN or
    # infinite,
    # so we print an info message instead of showing an empty legend.
    if show_legend:
        handles, labels = ax.get_legend_handles_labels()
        if handles:
            ax.legend()
        else:
            print(f'[INFO] Legend not shown for "{column_label}" because all
            # relevant values were NaN or infinite.')

def plot_box_with_stats_by_group(source_df,
                                column_name,
                                column_label,
                                group_by_column_name,

```

```

        group_by_column_label,
        grouped_df=None,
        show_legend=False):
fig, ax = plt.subplots(figsize=(10, 6))

if grouped_df is None:
    grouped_df = source_df.groupby([group_by_column_name])

# ***** Plot box chart *****
plot_box_by_group(source_df,
                  ax,
                  column_name,
                  column_label,
                  group_by_column_name,
                  group_by_column_label,
                  grouped_df,
                  show_legend)

plt.tight_layout()
plt.show()

# ***** Show statistical summary *****
print(f'{column_label} per {group_by_column_label}')
display(group_summary_stats(grouped_df, column_name))

def plot_histogram_all_and_by_group(source_df,
                                    x_column_name,
                                    x_column_label,
                                    group_by_column_name,
                                    group_by_column_label,
                                    count_label,
                                    grouped_df=None,
                                    as_percentage=False,
                                    include_overall=True):

    if include_overall:
        # ***** Plot histogram for overall data *****
        # print('\n')
        plt.figure(figsize=(12, 5))
        hist_overall = sns.histplot(source_df[x_column_name],
                                    bins=50,
                                    kde=True)

        plt.title(f'Overall Distribution of {x_column_label}')
        plt.xlabel(x_column_label)
        plt.ylabel(count_label)

        # Format x-axis as percentage (optional)
        if as_percentage:
            plt.gca().xaxis.set_major_formatter(PercentFormatter(xmax=1))

```

```

# Rotate x-axis labels
# plt.xticks(rotation=45)

# Add value labels on top of bars
for patch in hist_overall.patches:
    height = patch.get_height()
    if height > 0:
        plt.text(patch.get_x() + patch.get_width() / 2, height + 5,
↪f'{int(height)}',
                    ha='center', va='bottom', fontsize=8, rotation=45)

plt.tight_layout()
plt.show()

# ***** Show overall statistical summary *****
print(f'Overall {x_column_label}')
summary_df = pd.DataFrame([
    {
        'count': source_df[x_column_name].count(),
        'mean': source_df[x_column_name].mean(),
        'median': source_df[x_column_name].median(),
        'min': source_df[x_column_name].min(),
        'max': source_df[x_column_name].max()
    }
])
display(summary_df)
print('\n')

# ***** Plot histogram by group *****
plt.figure(figsize=(12, 5))
hist_by_group = sns.histplot(data=source_df, x=x_column_name,
↪hue=group_by_column_name,
                    bins=50, kde=True, element='bars',
↪stat='count', common_norm=False)
plt.title(f'Distribution of {x_column_label} by {group_by_column_label}')
plt.xlabel(x_column_label)
plt.ylabel(count_label)

if hist_by_group.get_legend() is not None:
    hist_by_group.get_legend().set_title(group_by_column_label)

# Format x-axis as percentage (optional)
if as_percentage:
    plt.gca().axis.set_major_formatter(PercentFormatter(xmax=1))

# Rotate x-axis labels

```

```

# plt.xticks(rotation=45)

# Add value labels per hue group
for container in hist_by_group.containers:
    for bar in container:
        height = bar.get_height()
        if height > 0:
            hist_by_group.annotate(f'{{int(height)}}', xy=(bar.get_x() + bar.
→get_width() / 2, height),
                                xytext=(0, 4), textcoords='offset_
→points', ha='center',
                                va='bottom', fontsize=8, rotation=45)

plt.tight_layout()
plt.show()

# ***** Show statistical summary by group *****
if grouped_df is None:
    grouped_df = source_df.groupby([group_by_column_name])

print(f'{{x_column_label}} by {{group_by_column_label}}')
group_summary_df = grouped_df.agg({x_column_name: ['count', 'mean',
→'median', 'min', 'max']})

# Flatten multi-level columns
group_summary_df.columns = [f'{{stat}}' for stat in group_summary_df.columns.
→get_level_values(1)]
group_summary_df = group_summary_df.reset_index()
display(group_summary_df)

def plot_scatter(source_df,
                 x_column_name,
                 x_column_label,
                 y_column_name,
                 y_column_label,
                 hue_column_name=None,
                 hue_column_label=None,
                 size_column=None):

    plt.figure(figsize=(12, 8))

    # Create scatter plot with optional hue and size
    scatter = sns.scatterplot(
        data=source_df,
        x=x_column_name,
        y=y_column_name,
        hue=hue_column_name,
        size=size_column,

```

```

        alpha=0.7
    )

    # Labels and title
    plt.xlabel(x_column_label)
    plt.ylabel(y_column_label)

    title = f'{x_column_label} vs {y_column_label}'
    if hue_column_label:
        title = f'{x_column_label} vs {y_column_label} by {hue_column_label}'
        plt.legend(title=hue_column_label)
    plt.title(title)

    plt.grid(True, linestyle='--', alpha=0.6)
    plt.show()

def plot_heatmap(source_df, column_name_label_map, title):
    plt.figure(figsize=(10, 8))

    # Derive column names
    column_names = list(column_name_label_map.keys())

    # Compute correlation and rename for display
    correlation = source_df[column_names].corr()
    correlation.rename(index=column_name_label_map,
        ↪ columns=column_name_label_map, inplace=True)

    sns.heatmap(correlation, annot=True, cmap='coolwarm', fmt='.2f')
    plt.title(title)
    plt.show()

```

#### 4.2.5 Preliminary Data Integrity Issues

Before conducting visual analysis, several inconsistencies were identified in the dataset. These issues affect key usage variables such as `drives`, `sessions`, and `driving_days`, and must be addressed to ensure accurate interpretation and reliable churn modeling.

```
[26]: summarize_waze_column_stats_multi(['driving_days', 'driven_km_drives'],
    ↪ 'drives', 'Drives')
```

<IPython.core.display.HTML object>

	total	non-null	null	min	max
driving_days	106	106	0	0.000000	28.00000
driven_km_drives	106	106	0	628.853609	16480.93908

<IPython.core.display.HTML object>

	total	non-null	null	min	max
driving_days	14	14	0	0.000000	17.00000
driven_km_drives	14	14	0	1410.581332	12316.76797

<IPython.core.display.HTML object>

	total	non-null	null	min	max
driving_days	88	88	0	0.000000	28.00000
driven_km_drives	88	88	0	628.853609	16480.93908

<IPython.core.display.HTML object>

	total	non-null	null	min	max
driving_days	4	4	0	4.00000	17.000000
driven_km_drives	4	4	0	1363.20614	5597.490052

**Drives = 0 with Distance Logged** This finding presents a logical contradiction. A user cannot accumulate hundreds or even thousands of kilometers of distance traveled (**driven\_km\_drives**) without completing any recorded trips (**drives**). This inconsistency suggests a failure in the trip logging mechanism, where the event counter (**drives**) did not register trips for 106 users while the distance tracker (**driven\_km\_drives**) continued to capture activity. As a result, this portion of the data is not reliable for straightforward analysis and must be cleaned or corrected before inclusion in any predictive modeling.

```
[27]: summarize_waze_column_stats_multi(['driven_km_drives', 'activity_days'],  
    ↪ 'sessions', 'Sessions')
```

<IPython.core.display.HTML object>

	total	non-null	null	min	max
driven_km_drives	105	105	0	628.853609	16480.93908
activity_days	105	105	0	0.000000	30.00000

<IPython.core.display.HTML object>

	total	non-null	null	min	max
driven_km_drives	14	14	0	1410.581332	12316.76797
activity_days	14	14	0	0.000000	19.00000

<IPython.core.display.HTML object>

	total	non-null	null	min	max
driven_km_drives	87	87	0	628.853609	16480.93908
activity_days	87	87	0	0.000000	30.00000

<IPython.core.display.HTML object>

	total	non-null	null	min	max
driven_km_drives	4	4	0	1363.20614	5597.490052
activity_days	4	4	0	9.00000	26.000000

**Sessions = 0 with Activity Recorded** This issue mirrors the previous finding. A user must initiate a session to log driving distance or activity days. It is highly improbable for a user to accumulate more than 600 kilometers and remain active for nearly an entire month without a single recorded session. This points to a flaw in the session recording process: sessions were in fact initiated and activity occurred, but the final session count was incorrectly stored as zero. The evidence confirms that these 105 users were active but were misrepresented by the `sessions` metric.

```
[28]: summarize_waze_column_stats_multi(['drives', 'driven_km_drives'],  
    ↪ 'driving_days', 'Driving days')
```

<IPython.core.display.HTML object>

	total	non-null	null	min	max
drives	1024	1024	0	0.000000	407.00000
driven_km_drives	1024	1024	0	159.444055	16321.74737

<IPython.core.display.HTML object>

	total	non-null	null	min	max
drives	393	393	0	0.000000	382.00000
driven_km_drives	393	393	0	195.996535	14936.04257

<IPython.core.display.HTML object>

	total	non-null	null	min	max
drives	590	590	0	0.000000	407.00000
driven_km_drives	590	590	0	159.444055	16321.74737

<IPython.core.display.HTML object>

	total	non-null	null	min	max
drives	41	41	0	2.000000	224.00000
driven_km_drives	41	41	0	799.048257	12961.44177

**Driving Days = 0 with Trips Logged** This is the most significant and widespread data flaw. It is logically impossible for a user to complete hundreds of trips and drive thousands of kilometers without any recorded driving days. This error affects over 1,000 users across both retained and churned groups. Since `driving_days` was previously identified as a strong predictor of churn (with retained users averaging 14 days vs. 6 for churned users), the unreliability of this metric undermines a key variable for prediction.



**Overall Conclusion** These three findings confirm that essential usage metrics—`drives`, `sessions`, and `driving_days`—are compromised by systematic logging errors. Left uncorrected, they will distort both descriptive analysis and predictive modeling. To address this, zero values that conflict with positive distance (`driven_km_drives`) should either be replaced with `NaN` to exclude them from calculations or adjusted to corrected values inferred from related activity. Cleaning these anomalies is a prerequisite for producing a reliable churn prediction model.

**PACE: Analyze Stage Questions** **Question 1:** Will the available information be sufficient to achieve the goal based on your intuition and the analysis of the variables?

**Question 2:** How would you build summary dataframe statistics and assess the min and max range of the data?

**Question 3:** Do the averages of any of the data variables look unusual? Can you describe the interval data?

**Answer 1:** > Yes, the dataset contains enough information to build a baseline churn prediction model. However, several integrity issues were identified, such as users with positive kilometers but zero recorded `drives`, `sessions`, or `driving_days`. These anomalies must be addressed through cleaning or validation before modeling. While the available variables are sufficient for initial work, additional historical (multi-month) data would provide deeper behavioral context and improve model reliability.

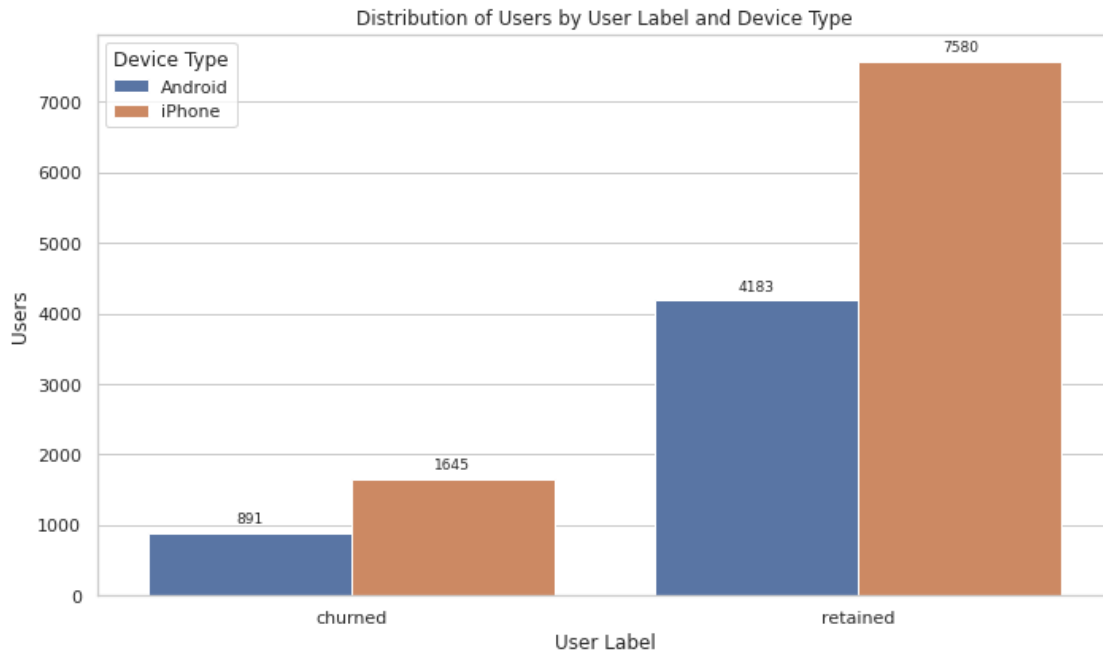
**Answer 2:** > We would use the `.describe()` method in Pandas. The minimum and maximum values in its output directly show the range for each numerical variable. This allows quick detection of outliers and validation of whether values fall within logical limits.

**Answer 3:** > The averages themselves are not problematic, but they reveal skewness due to extreme outliers. For example, maximum values for `sessions`, `drives`, and `driven_km_drives` are far higher than their means, suggesting the presence of heavy users. In addition, some averages may be misleading if underlying variables contain contradictions (*e.g., users with kilometers logged but zero drives*). The numerical variables in this dataset represent ratio data, not interval data. Ratio data has a meaningful zero point (*e.g., zero kilometers means no distance traveled*) and allows for meaningful comparisons using ratios (*e.g., 10 km is twice 5 km*).

#### 4.2.6 Visualization

##### Comparison of Churned vs. Retained Users by Device Type

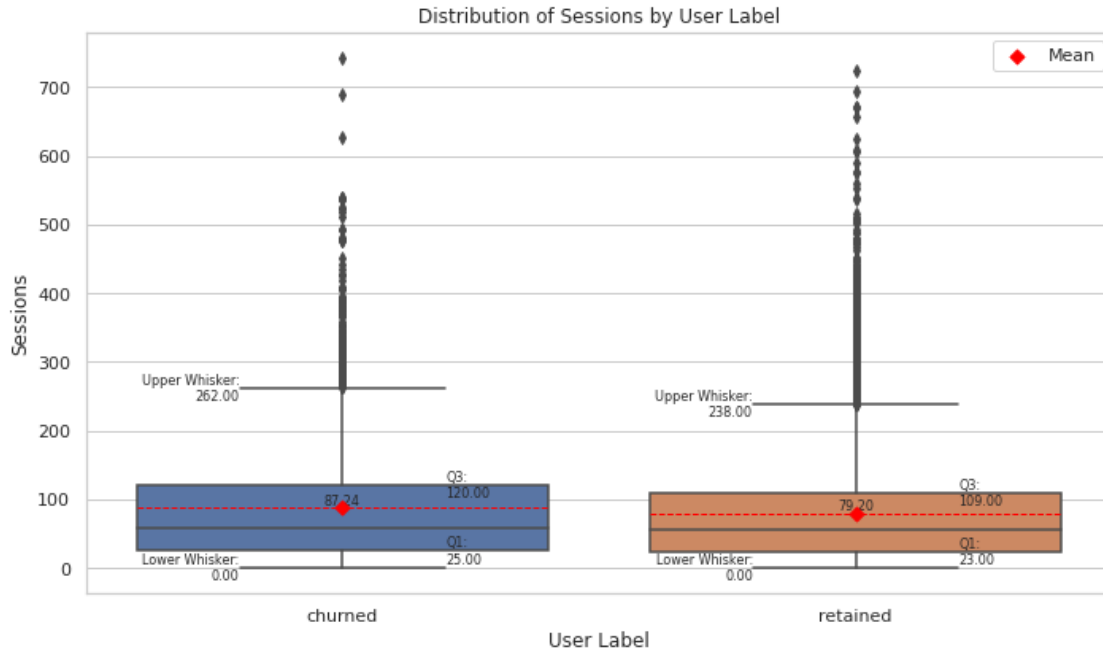
```
[29]: plot_bar_by_group(waze_users_df, 'label', 'User Label', 'device', 'Device_  
      ↪Type', 'Users')
```



This chart compares churned and retained users across iPhone and Android devices. The proportions are nearly identical in both groups: approximately 65% iPhone and 35% Android. This indicates that device type is not a meaningful driver of churn, as churn occurs at similar rates across platforms.

#### Distribution of User Sessions for Churned vs. Retained Groups

```
[30]: plot_box_with_stats_by_group(waze_users_df, 'sessions', 'Sessions', 'label',  
    ↪ 'User Label', waze_users_label_group, True)
```



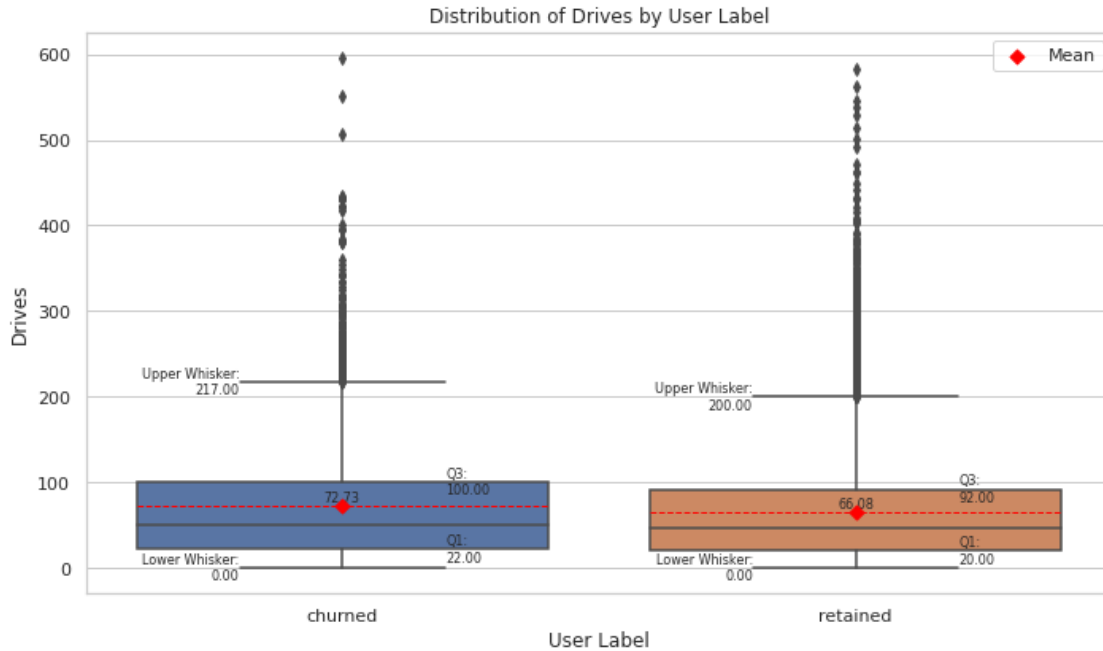
#### Sessions per User Label

	label	mean	median	Q1	Q3	IQR	min	max
0	churned	87.24	59.0	25.0	120.0	95.0	0	743
1	retained	79.20	56.0	23.0	109.0	86.0	0	725

The distribution of sessions reveals that both churned and retained users have similar medians, but churned users exhibit slightly higher variability. The presence of significant outliers among churned users suggests that while most behave similarly to retained users, a small subset engages in an unusually high number of sessions.

#### Distribution of User Drives for Churned vs. Retained Groups

```
[31]: plot_box_with_stats_by_group(waze_users_df, 'drives', 'Drives', 'label', 'User_
      ↪Label', waze_users_label_group, True)
```



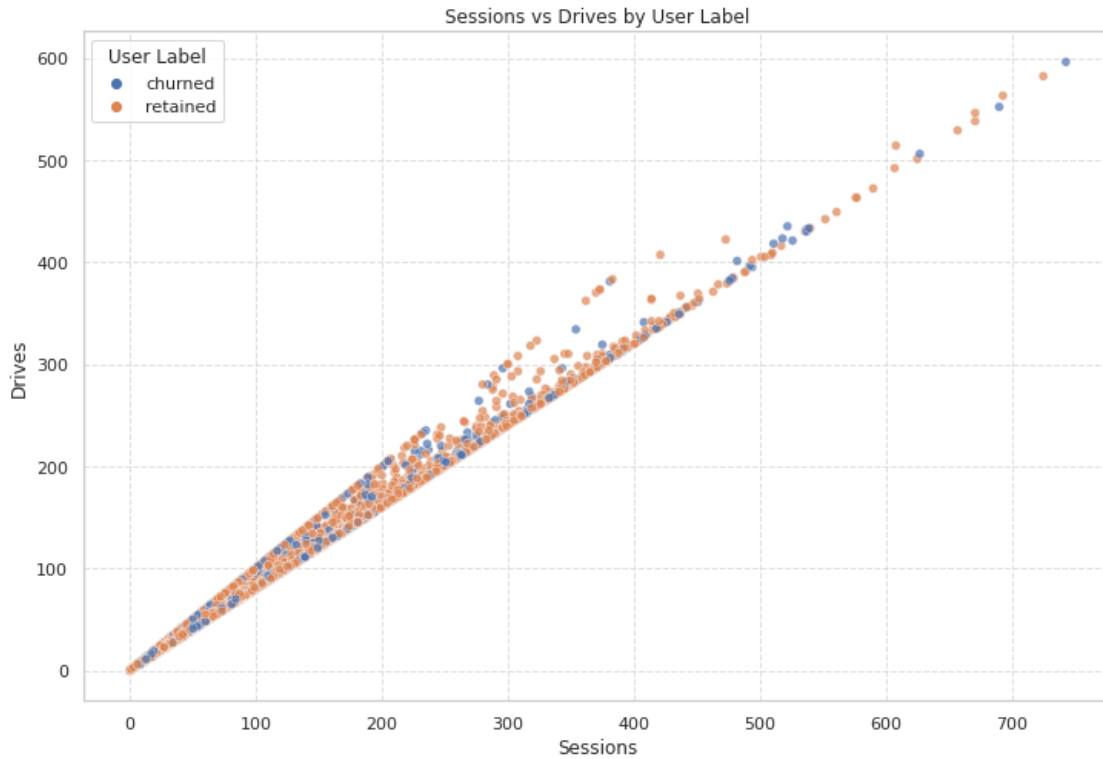
#### Drives per User Label

	label	mean	median	Q1	Q3	IQR	min	max
0	churned	72.73	50.0	22.0	100.0	78.0	0	596
1	retained	66.08	47.0	20.0	92.0	72.0	0	582

The boxplot indicates that churned users typically have more drives than retained users, with medians of 50 and 47, respectively. However, churned users also display greater variability and a higher concentration of outliers. This suggests that churned users may exhibit more intense but irregular usage patterns.

#### Relationship Between Sessions and Drives by User Retention Status


```
[32]: plot_scatter(waze_users_df, 'sessions', 'Sessions', 'drives', 'Drives',
    ↪ 'label', 'User Label')
```



The scatter plot highlights a strong positive correlation between sessions and drives for both groups, confirming that more sessions generally translate into more drives. However, churned users appear slightly more concentrated at higher drive counts, suggesting they may complete more trips within a comparable number of sessions.

#### Relationship Between Drives and Kilometers Driven by User Retention Status


```
[33]: plot_scatter(waze_users_df, 'drives', 'Drives', 'driven_km_drives', 'Kilometers_Driven', 'label', 'User Label')
```



This visualization demonstrates a clear positive relationship between the number of drives and total kilometers driven. While both churned and retained users follow the same overall trend, churned users show a tendency toward longer total distances, implying that they may use Waze for longer or more frequent trips.

#### Relationship Between Sessions and Kilometers Driven by User Retention Status


```
[34]: plot_scatter(waze_users_df, 'sessions', 'Sessions', 'driven_km_drives',  
                 ↪ 'Kilometers Driven', 'label', 'User Label')
```



Sessions correlate strongly with kilometers driven, but churned users display higher kilometer counts at comparable session levels. This suggests that churned users engage in longer trips per session, reinforcing the idea that their driving behavior differs in intensity compared to retained users.

### Comparison of Driving Days per Month for Churned vs. Retained Users

```
[35]: plot_histogram_all_and_by_group(waze_users_df, 'driving_days', 'Driving Days',
                                     'label', 'User Label', 'Users',
                                     ↪waze_users_label_group, include_overall=False)
```




### Driving Days by User Label

	label	count	mean	median	min	max
0	churned	2536	7.218060	6.0	0	29
1	retained	11763	13.252827	14.0	0	30

The histogram shows a notable separation between groups. Retained users drive on more days per month (median ~14 days), while churned users drive on fewer (median ~6 days). This finding indicates that higher frequency of app engagement across the month is a key differentiator of retention.

### Comparison of Activity Days per Month for Churned vs. Retained Users

```
[36]: plot_histogram_all_and_by_group(waze_users_df, 'activity_days', 'Activity Days',  
                                     'label', 'User Label', 'Users',  
                                     ↪waze_users_label_group, include_overall=False)
```



### Activity Days by User Label


	label	count	mean	median	min	max
0	churned	2536	9.644716	8.0	0	31
1	retained	11763	16.816628	17.0	0	31

Similar to driving days, retained users have significantly more app activity days (median ~17) compared to churned users (median ~8). This reinforces the pattern that consistent and sustained engagement with the app is strongly associated with user retention.

### Correlation Matrix of Key User Activity Variables



```
[37]: plot_heatmap(waze_users_df,
                  {
                      'sessions': 'Sessions', 'drives': 'Drives',
                      'n_days_after_onboarding': 'Days After Onboarding',
                      'activity_days': 'Activity Days', 'driving_days': 'Driving Days'
                  },
                  'Correlation Matrix of Numerical Variables')
```



The correlation heatmap reveals strong positive relationships among activity-related variables such as sessions, drives, activity days, and driving days, reflecting their shared role in measuring user engagement. By contrast, days after onboarding shows a weaker relationship with current activity, suggesting that longevity alone does not predict engagement or churn.

### 4.3 PACE: Construct

**Note:** The Construct stage does not apply to this workflow. The PACE framework can be adapted to fit the specific requirements of any project.

### 4.4 PACE: Execute

Consider the questions in your PACE Strategy Document and those below to craft your response:

**Question 1:** Given your current knowledge of the data, what would you initially recommend to your manager to investigate further prior to performing exploratory data analysis?

**Question 2:** What data initially presents as containing anomalies?

**Question 3:** What additional types of data could strengthen this dataset?

**Answer 1:** > We recommend first investigating the significant inconsistencies found in the data. Specifically, many users show a positive value for `driven_km_drives` but have **zero drives** or **driving\_days**. This is a critical data quality issue that must be addressed before any analysis can be trusted.

In addition, the **700 missing values** in the `label` column also need to be investigated. Since `label` is our target variable for predicting churn, understanding why these values are missing is essential. They may represent new users who have not yet been classified, or they could be the result of data collection errors. Knowing the reason behind these missing values will determine whether these rows should be removed, imputed, or treated as a distinct group. Resolving both the inconsistencies and the missing labels will provide a stronger foundation for reliable analysis and modeling.

**Answer 2:** > There are outliers in columns such as `sessions`, `drives`, `total_sessions`, `driven_km_drives`, and `duration_minutes_drives` that could skew the analysis and model performance. In addition, some columns show inconsistencies. For example, a few users have zero `drives` or `sessions` but still show positive driving days or kilometers driven, while over 1,000 users have significant drives and kilometers but zero driving days. These are not just outliers, since these fields are expected to align, the mismatches suggest potential data errors that need to be validated and cleaned before modeling.

**Answer 3:**

> 1. Having historical data of users, rather than just a single monthly snapshot, would likely improve the prediction accuracy of our machine learning model. It would let us see trends in user activity leading up to churn, like a gradual decrease in sessions or drives. This helps us understand the “why” much better than a single point in time. 2. Having data on app and feature usage would also be valuable. Beyond just monthly totals, more granular data—such as the frequency of app use within a day or week, or engagement with specific Waze features—could provide deeper understanding. For example, user interactions like actively reporting traffic, using the carpool feature, connecting with friends, or submitting bug reports or contacting customer support (a potential indicator of frustration leading to churn) can reveal important behavioral patterns. 3. Having user profile data—such as age, general location, or the types of routes they typically take (e.g., daily commutes vs. long-distance travel)—could also provide valuable context for understanding user behavior.

#### 4.4.1 Task 3. Conclusion

Recall that your supervisor, May Santer, asked you to share your findings with the data team in an executive summary. Consider the following questions as you prepare to write your summary. Think about key points you may want to share with the team, and what information is most relevant to the user churn project.

##### Questions:

1. Did the data contain any missing values? How many, and which variables were affected? Was there a pattern to the missing data?
  - > Yes, the data had missing values. The `label` column, which tells us if a user has churned or retained, had **700 missing values**. Earlier analysis show that there didn't seem to be a strong pattern to the missing data. The average user behavior for those with a missing label was very similar to the rest of the dataset.
  - In addition to missing labels, we also found significant inconsistencies across related columns. For example, some users recorded positive kilometers driven (`driven_km_drives`) despite having zero `drives` or `driving_days`, while over 1,000 users had substantial `drives` and `driven_km_drives` but zero `driving_days`. Since these fields are supposed to align, these mismatches suggest potential data quality issues that will need to be validated and cleaned before modeling.
2. What is a benefit of using the median value of a sample instead of the mean?
  - > A big benefit of using the median is that it's not affected by outliers. The mean can get skewed pretty easily by extremely high or low values. For example, if one user drove 20,000 km in a month, that would pull the average (mean) way up. The median, on the other hand, would stay closer to the middle of the pack, giving us a more representative idea of a "typical" user's behavior.
3. Did your investigation give rise to further questions that you would like to explore or ask the Waze team about?
  - > Yes, these questions are:
    - **Why is the `label` missing for 700 users?**  
We found that their behavior is pretty standard, so it's not because they're new or inactive. We need to know if there's a technical reason for this missing data or if these users were intentionally left out.
    - **Why are `km_per_driving_day` and `drives_per_driving_day` so high for churned users?**  
The median for churned users is significantly higher than for retained users. This is a very interesting pattern. Are these users driving long distances for work (e.g., couriers) and perhaps find another app more suitable, or are they a different type of user altogether? We need to dig into this to understand the "why."
    - **What's the relationship between sessions and `total_sessions`?**  
The descriptive stats show that some users have a low number of sessions this month but a very high total number of sessions. This suggests they might be long-time users who are becoming less active. This could be a very strong predictor for churn.
    - **Why do some users record distance and trips without corresponding sessions,**

### drives, or driving days?

Several inconsistencies were identified where users accumulated kilometers or trips while the related counters remained at zero. These contradictions point to possible data logging failures. It would be critical to know whether these errors stem from system design, technical glitches, or post-processing steps, since they impact key variables used in churn modeling.

4. What percentage of the users in the dataset were Android users and what percentage were iPhone users?

> The percentages for the entire dataset are very close to the percentages of those with a null label, as well as the percentages within the churned and retained groups.

#### Overall:

iPhone users: **64.48%** | Android users: **35.52%**

#### Null label:

iPhone users: **63.86%** | Android users: **36.14%**

#### Churned users:

iPhone users: **64.87%** | Android users: **35.13%**

#### Retained users:

iPhone users: **64.44%** | Android users: **35.56%**

5. What were some distinguishing characteristics of users who churned vs. users who were retained?

> **Churned users** have a much higher median `km_per_driving_day` (about 698 km) compared to retained users (about 290 km). Similarly, they have a much higher median `drives_per_driving_day` (10 drives) compared to retained users (4 drives). This suggests that users who churn may be more intense, frequent drivers on the days they're using the app.

**Retained users** on the other hand, have a higher median `n_days_after_onboarding` (1843 days) compared to churned users (1321 days), which hints that they are generally longer-term users. They also have more `activity_days` and `driving_days` during the month.

6. Was there an appreciable difference in churn rate between iPhone users vs. Android users?

> No, there was not. The percentage of iPhone users and Android users within the churned group (64.87% and 35.13% respectively) is almost identical to their proportion in the entire dataset and the retained group. This means that a user's device type doesn't appear to be a strong predictor of churn.

#### Overall:

The dataset provides a strong basis for churn analysis, but two major data quality issues stand out: the 700 missing labels and the inconsistencies across related activity columns. Both need to be resolved to ensure reliable results.

---

```
[38]: waze_users_df.head(10)
```

```
[38]:   id  label  sessions  drives  total_sessions  n_days_after_onboarding  \
0   0  retained      283     226      296.748273                2276
1   1  retained      133     107      326.896596                1225
2   2  retained      114      95      135.522926                2651
3   3  retained       49      40       67.589221                 15
```

4	4	retained	84	68	168.247020	1562
5	5	retained	113	103	279.544437	2637
6	6	retained	3	2	236.725314	360
7	7	retained	39	35	176.072845	2999
8	8	retained	57	46	183.532018	424
9	9	churned	84	68	244.802115	2997

	total_navigations_fav1	total_navigations_fav2	driven_km_drives	\
0	208	0	2628.845068	
1	19	64	13715.920550	
2	0	0	3059.148818	
3	322	7	913.591123	
4	166	5	3950.202008	
5	0	0	901.238699	
6	185	18	5249.172828	
7	0	0	7892.052468	
8	0	26	2651.709764	
9	72	0	6043.460295	

	duration_minutes_drives	activity_days	driving_days	device	\
0	1985.775061	28	19	Android	
1	3160.472914	13	11	iPhone	
2	1610.735904	14	8	Android	
3	587.196542	7	3	iPhone	
4	1219.555924	27	18	Android	
5	439.101397	15	11	iPhone	
6	726.577205	28	23	iPhone	
7	2466.981741	22	20	iPhone	
8	1594.342984	25	20	Android	
9	2341.838528	7	3	iPhone	

	km_per_drive	km_per_driving_day	drives_per_driving_day
0	11.632058	138.360267	11.894737
1	128.186173	1246.901868	9.727273
2	32.201567	382.393602	11.875000
3	22.839778	304.530374	13.333333
4	58.091206	219.455667	3.777778
5	8.749890	81.930791	9.363636
6	2624.586414	228.224906	0.086957
7	225.487213	394.602623	1.750000
8	57.645864	132.585488	2.300000
9	88.874416	2014.486765	22.666667

```
[39]: waze_users_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
```

Data columns (total 16 columns):

#	Column	Non-Null Count	Dtype
0	id	14999 non-null	int64
1	label	14299 non-null	category
2	sessions	14999 non-null	int64
3	drives	14999 non-null	int64
4	total_sessions	14999 non-null	float64
5	n_days_after_onboarding	14999 non-null	int64
6	total_navigations_fav1	14999 non-null	int64
7	total_navigations_fav2	14999 non-null	int64
8	driven_km_drives	14999 non-null	float64
9	duration_minutes_drives	14999 non-null	float64
10	activity_days	14999 non-null	int64
11	driving_days	14999 non-null	int64
12	device	14999 non-null	category
13	km_per_drive	14999 non-null	float64
14	km_per_driving_day	14999 non-null	float64
15	drives_per_driving_day	14992 non-null	float64

dtypes: category(2), float64(6), int64(8)

memory usage: 1.6 MB

```
[40]: waze_users_df.describe()
```

```
[40]:
```

	id	sessions	drives	total_sessions	\
count	14999.000000	14999.000000	14999.000000	14999.000000	
mean	7499.000000	80.633776	67.281152	189.964447	
std	4329.982679	80.699065	65.913872	136.405128	
min	0.000000	0.000000	0.000000	0.220211	
25%	3749.500000	23.000000	20.000000	90.661156	
50%	7499.000000	56.000000	48.000000	159.568115	
75%	11248.500000	112.000000	93.000000	254.192341	
max	14998.000000	743.000000	596.000000	1216.154633	

	n_days_after_onboarding	total_navigations_fav1	\
count	14999.000000	14999.000000	
mean	1749.837789	121.605974	
std	1008.513876	148.121544	
min	4.000000	0.000000	
25%	878.000000	9.000000	
50%	1741.000000	71.000000	
75%	2623.500000	178.000000	
max	3500.000000	1236.000000	

	total_navigations_fav2	driven_km_drives	duration_minutes_drives	\
count	14999.000000	14999.000000	14999.000000	
mean	29.672512	4039.340921	1860.976012	

std	45.394651	2502.149334	1446.702288
min	0.000000	60.441250	18.282082
25%	0.000000	2212.600607	835.996260
50%	9.000000	3493.858085	1478.249859
75%	43.000000	5289.861262	2464.362632
max	415.000000	21183.401890	15851.727160

	activity_days	driving_days	km_per_drive	km_per_driving_day \
count	14999.000000	14999.000000	1.499900e+04	1.499900e+04
mean	15.537102	12.179879	inf	inf
std	9.004655	7.824036	NaN	NaN
min	0.000000	0.000000	1.008775e+00	3.022063e+00
25%	8.000000	5.000000	3.323065e+01	1.672804e+02
50%	16.000000	12.000000	7.488006e+01	3.231459e+02
75%	23.000000	19.000000	1.854667e+02	7.579257e+02
max	31.000000	30.000000	inf	inf

	drives_per_driving_day
count	1.499200e+04
mean	inf
std	NaN
min	0.000000e+00
25%	1.800000e+00
50%	4.666667e+00
75%	1.216667e+01
max	inf

#### 4.4.2 Data Analyst Notes

**inf and NaN Values of Derived Columns** In the resulting DataFrame after my initial cleaning, I've added three computed columns to aid in my analysis: `* km_per_drive = driven_km_drives / drives` `* km_per_driving_day = driven_km_drives / driving_days` `* drives_per_driving_day = drives / driving_days`

Using `pandas`' `describe()`, the values for `mean`, standard deviation (`std`), and `max` are `inf` and `NaN` for these computed columns. It was observed, that in `drives` and `driving_days` there are zero values (0.0) present. These columns are being used as divisors for our computed columns. In `pandas`, when a floating-point number is divided by zero, the result is positive or negative infinity (`inf`). The `NaN` values for the standard deviation (`std`) in these columns are a direct consequence of the `inf` values. It's not possible to compute the standard deviation when a dataset contains infinite values.

#### inf Explanation

**Python Behavior** In standard Python, dividing by zero raises a `ZeroDivisionError`. This is because the result is mathematically undefined and cannot be represented as a finite number. This

forces you to handle the error explicitly instead of allowing invalid results to continue.

**NumPy and Pandas Behavior** In contrast, when working with NumPy arrays or Pandas DataFrames that store values as floating-point numbers (float64), division follows the **IEEE 754 standard** for floating-point arithmetic. This standard defines special values for division by zero:

- Positive number / 0.0 → positive infinity (**inf**)
- Negative number / 0.0 → negative infinity (**-inf**)
- Zero / 0.0 → “Not a Number” (**NaN**).

When calculating the derived ratio columns in `waze_users_df`, any division by zero results in **inf** or **NaN** values instead of raising an error. This allows Pandas to compute the entire column without interruption, but it also means that these special values can appear in the results. In turn, statistical summaries such as `mean` or `std` may return **NaN** if the column contains **inf** values, because meaningful variance calculations cannot be performed with infinite values.

```
[41]: # This will raise a ZeroDivisionError
try:
    py_result = 50.0 / 0.0
except ZeroDivisionError as error:
    print(f'Python behavior: Error: {error}')

# This will NOT raise an error and will return inf
numpy_result = np.divide(50.0, 0.0)
print(f'NumPy behavior: {numpy_result}')

# This is what happens inside a Pandas DataFrame column
numerator_se = pd.Series([50.0, 10.0])
denominator_se = pd.Series([0.0, 2.0])
result_se = numerator_se / denominator_se
print(f'Result:\n{result_se}')
```

Python behavior: Error: float division by zero

NumPy behavior: inf

Result:

0     inf

1     5.0

dtype: float64

**inf values in max** The maximum value appears as **inf** because, under Python’s comparison rules and the IEEE 754 floating-point standard, positive infinity is greater than any finite number. When `pandas.describe()` calculates a column’s maximum, it compares each value to the current highest; since **inf** is always larger than any finite value, it becomes the reported maximum.

**NaN Explanation** The standard deviation becomes **NaN** when a column contains **inf** values because the calculation cannot be performed if the `mean` is infinite. In statistics, standard deviation measures the “spread around a mean”, meaning how far the values vary from the average. A small



standard deviation means values are close to the mean, while a large one means they are more scattered.

```
[42]: # Create a boolean mask to get inf values in km_per_drive column
km_per_drive_inf = np.isinf(waze_users_df['km_per_drive'])

# Apply the mask to our DataFrame
km_per_drive_inf_df = waze_users_df[km_per_drive_inf]

km_per_drive_inf_df[['driven_km_drives', 'drives', 'km_per_drive']].head(5)
```

```
[42]:      driven_km_drives  drives  km_per_drive
25          5702.339466      0          inf
97          6668.844350      0          inf
217         6103.881670      0          inf
339         2520.850896      0          inf
485         1363.206140      0          inf
```

```
[43]: km_per_drive_df = waze_users_df[['km_per_drive']]
print(f"Mean\n{km_per_drive_df.mean(numeric_only=True)}")
print(f"\nStandard Deviation\n{km_per_drive_df.std(numeric_only=True)}")
```

```
Mean
km_per_drive    inf
dtype: float64
```

```
Standard Deviation
km_per_drive    NaN
dtype: float64
```

```
[44]: # Demonstrating the result by creating a very small sample
data_with_inf = np.array([10, 20, 30, float('inf')])

# Calculate the mean
mean_value = np.mean(data_with_inf)
print(f"The mean of the array is: {mean_value}")

# Calculate the standard deviation
std_value = np.std(data_with_inf)
print(f"The standard deviation of the array is: {std_value}")
```

```
The mean of the array is: inf
```

```
The standard deviation of the array is: nan
```

The standard deviation is calculated using the following formula:

$$\sigma = \sqrt{\frac{\sum (x_i - \mu)^2}{N}}$$

- $x_i$  is each individual value in the dataset.
- $\bar{x}$  is the mean of the dataset.
- $N$  is the number of values.

**Given observations:** [10, 20, 30, inf]

This formula fails when your data contains **inf**:

1. **Calculate the Mean ( $\bar{x}$ ):** The mean is the sum of all values divided by the count. When you add a finite number to **inf**, the result is **inf**.

```
mean = (10 + 20 + 30 + inf) / 4
      = inf / 4
      = inf
```

2. **Calculate the Deviation from the Mean ( $x_i - \bar{x}$ ):**

```
-inf = 10 - inf
-inf = 20 - inf
-inf = 30 - inf
NaN = inf - inf
```

3. **Calculate the Sum of Squared Deviations ( $(x_i - \bar{x})^2$ ):**

```
inf = (-inf)^2
NaN = NaN^2
```

```
NaN = (inf + inf + inf + NaN)
```

Since the numerator of the formula becomes **NaN**, the final result of the standard deviation calculation is also **NaN**.

**Mitigation Strategies** Division by zero in derived columns can produce **inf** or **NaN** values, which may distort statistical summaries. Common approaches include:

- Filtering out rows where the denominator is zero, removing them entirely from the dataset.
- Replacing **inf** and **NaN** with a suitable value (such as 0).
- Using **NaN** for invalid results to retain the rows while ensuring they are excluded from calculations.

These steps help maintain the accuracy of summary statistics and visualizations.

**Best practice:** For most analyses, using **NaN** is recommended. It preserves data integrity, keeps the dataset complete, and ensures statistical functions automatically ignore invalid values.

---

## Why Our Most Active Users Might Be Leaving

The churned group racks up more sessions and drives than those who stay — and it may not be a coincidence.

Analysis of the available data shows that churned users have, on average, higher sessions, drives, and related activity metrics compared to retained users. One possible explanation — though not directly verifiable with the current dataset — is that heavy usage of Waze could lead to increased battery drain and data consumption, which might prompt some users to switch to alternative navigation solutions. However, because the dataset does not include direct measures of device battery usage, data consumption, or user feedback, this interpretation remains a hypothesis rather than a confirmed cause of churn.

---

**Congratulations!** You’ve completed this lab. However, you may not notice a green check mark next to this item on Coursera’s platform. Please continue your progress regardless of the check mark. Just click on the “save” icon at the top of this notebook to ensure your work has been logged.